

Atomic Consistency Memory in BAMP systems

Yosef Goren

On '**Atomic Read/Write Memory in Signature-Free Byzantine Asynchronous Message-Passing Systems**'.

A paper by:

Achour Mostefaoui, Matoula Petrolia, Michel Raynal, Claude Jard

Table of Contents

1 Introduction

- Background
- System Model
- Specifications

2 Algorithm

- Non-Atomic Algorithm
- Algorithm for failure-free systems
- BAMS Algorithm (from the paper)

3 Analysis

- Termination Properties
- Atomicity Properties
- Piecing it all together

4 Conclusions

- What we have seen
- Further Work

Why care about implementing registers?

What we get

This implementation provides a reduction from Message Passing models to Atomic Consistency Memory models.

What it can be used for

Many distributed algorithms are based on atomic memory; this reduction provides instant implementations of these algorithms in message passing systems.

Examples

- Atomic, multi-writer multi-reader registers
- Concurrent time-stamp systems
- Atomic snapshot scan

Previous Atomic Register Algorithms

BAMP: Byzantine Asynchronous Message Passing

A distributed system of n processes p_1, p_2, \dots, p_n .

Byzantine

A byzantine process is one that acts arbitrarily, it may crash or even send 'malicious' messages to correct processes.

Let t be the number of byzantine processes, we assume $t < \frac{n}{3}$.

Asynchronous

A message sent from p_i to p_j may take any amount of time to arrive.

Signature Free

Many algorithms cope with byzantine processes by requiring them to sign messages, thus requiring assuming cryptographic primitives to be correct, it is not the case here.

Reliable Broadcast Abstraction

Based on a separate paper, we can use a reliable broadcast algorithm as a 'black box' (in *BAMP* systems).

Guarantees

The reliable broadcast will have syntax '*r_broadcast m*', and guarantees that the message *m* arrives at all correct processes eventually.

Single Writer Multiple Reader Registers (*SWMR*)

A single process can write; everyone can read.

Single Writer & Byzantine Processes

If all shared memory can be written by all processes - a single Byzantine process can destroy it.

Local Copies

Each process p_i has Reg_i , but can only write to $Reg_i[i]$.

p_1	p_2	p_3
$Reg_1[1]$	$Reg_2[1]$	$Reg_3[1]$
$Reg_1[2]$	$Reg_2[2]$	$Reg_3[2]$
$Reg_1[3]$	$Reg_2[3]$	$Reg_3[3]$

Atomic Consistency

Atomic Consistency requires no concurrent actions to be interleaved, it is also known as **Linearizability**.

(Almost) Formal Definition

'for any execution of the system, there is some way of totally ordering the reads and writes so that the values returned by the reads are the same as if the operations had been performed in that order, with no overlapping.'

- 'On Interprocess Communication', Lamport (1985).

Notations

We define these notations for any correct processes p_i, p_j :

Reads

$read[i, j, x]$ will refer to an invocation by p_i , to read $Reg_i[j]$ which returns the x 'th value written by p_j .

Writes

$write[i, y]$ will refer to the y 'th invocation by p_i , to write $Reg_i[i]$.

Termination Requirement

Let p_i be a correct process.

Write Termination

Each invocation of $Reg_i[i].write()$ terminates.

Read Termination

For any j , all invocations $Reg_i[j].read()$ terminate.

Consistency Requirement

Let p_i, p_j be correct processes, and p_k be (possibly) byzantine.

Associated sequence

We can associate a sequence $H_k[x]$ with p_k , s.t. if p_k is correct, $H_k[x]$ is the value written by $write[k, x]$.

Read followed by Write

if $read[j, i, x]$ terminates before $read[i, j, y]$ then $x \leq y$.

Write followed by Read

if $write[j, x]$ terminates before $read[i, j, y]$ then $x \leq y$.

No Read inversion

if $read[i, k, x]$ terminates before $read[j, k, y]$ starts then $x \leq y$.

Local variables

Each process has local variables reg, rsn - arrays of n , and an additional counter wsn .

Initialization

$$reg \leftarrow [(init_0, 0), \dots, (init_n, 0)].$$
$$rsn \leftarrow [0, \dots, 0].$$
$$wsn \leftarrow 0.$$

Write

```
Reg[i].write(v):  
    wsn := wsn + 1  
    r_broadcast WRITE(v, wsn)  
    wait WRITE_DONE(wsn) received from (n-t) processes  
    return
```

Read

```

Reg[j].read():
  rsn[j] := rsn[j] + 1
  r_broadcast READ(v, rsn[j])
  wait untill reg[j].sn is greater than the arguments of
    n-t messages STATE(rsn[j], _) from different processes.
  let (w, wsn) := reg[j]
  r_broadcast CATCH_UP(w, wsn)
  wait CATCH_UP_DONE(j, wsn) recived from (n-t) different processes.
  return w.
  
```

Message Handling

```

on reciving WRITE(v, wsn) from process j:
    wait (wsn = reg[j].sn + 1)
    reg[j] := (v, wsn)
    send WRITE_DONE(wsn) to process k

on reciving READ(v, rsn) from process k:
    send STATE(rsn, reg[j].sn) to process k

on reciving CATCH_UP(j, wsn) from process k:
    wait (reg[j].sn >= wsn)
    send CATCH_UP_DONE(j, wsn) to process k
  
```

Lemma 1

Lemma

If a correct process p_i receives a message m from a $r_broadcast(m)$ by another correct process - any other correct process will receive m .

Proof.

Immediate from the guarantees of the broadcast algorithm. □

Lemma 2

Lemma

Any two sets of processes of size $(n - t)$ must have at least one correct process in common.

Lemma 2 - Proof

Proof.

Denote the set of processes with P , and the set of faulty ones F .
Let $Q_1, Q_2 \subseteq P$ s.t. $|Q_1| = |Q_2| = n - t$.

$$\begin{aligned} |\overline{Q_1} \cup \overline{Q_2}| &\leq |\overline{Q_1}| + |\overline{Q_2}| \Rightarrow n - |\overline{Q_1} \cup \overline{Q_2}| \geq n - |\overline{Q_1}| - |\overline{Q_2}| \\ &\Rightarrow |\overline{\overline{Q_1} \cup \overline{Q_2}}| \geq n - t - t \\ &\Rightarrow |Q_1 \cap Q_2| \geq n - 2t > 3t - 2t = t = |F| \\ &\Rightarrow \exists p \in Q_1 \cap Q_2 \notin F \end{aligned}$$



Lemma 3

Lemma

Let p_i be a correct process. Any invocation of $\text{Reg}[i].\text{write}()$ terminates.

Proof.

When p_i invokes $\text{Reg}[i].\text{write}()$ it sends *WRITE* to all others, due to reliable broadcast - $n - t$ correct processes receive and handle this message eventually, thus send back *WRITE_DONE*.

At some point these arrive and $\text{Reg}[i].\text{write}()$ terminates. □

Lemma 4

Lemma

Lemma 5

Lemma 6

Lemma

Let p_i, p_j be two correct processes.

If $\text{read}[i, j, x]$ terminates before $\text{write}[j, y]$ starts, then $x < y$.

Lemma 7

Lemma

Let p_i, p_j be two correct processes.

If $\text{write}[i, x]$ terminates before $\text{read}[j, i, y]$ starts, then $x \leq y$.

Lemma 8

Lemma

Let p_i, p_j be two correct processes.

If $\text{read}[i, k, x]$ terminates before $\text{read}[j, k, y]$ starts, then $x \leq y$.

Theorem

The algorithm showcased implements an array of n SWMR registers with atomic consistency, in BAMP with $t < \frac{n}{3}$ systems.

Proof.

We have seen required termination properties in lemmas 3,4 and atomicity properties in lemmas 5,6,7,8.



What we have seen

Taxonomy and building blocks

Atomic Consistency, SWMR, Reliable Broadcast

Shared Memory Algorithms

We have seen some naive ideas for sharing memory, and a correct algorithm for *BAMP* systems.

Correctness Proof

Each of the algorithm's wanted properties has been shown.

Sequential Consistency too much?

Runtime Limitations

Requiring a system to implement Atomic Consistency is a very strong requirement and often comes at a steep runtime cost.

Alternative Models: $AC \subseteq SC \subseteq RC$

Is an algorithm for (only) Sequential Consistency possible?
Or better yet - an algorithm for Release Consistency with some sort of '*fence*' operation?

Exploding Serial Numbers

Number of messages sent is unbounded, memory complexity is logarithmic with number of messages sent (due to counters).

Reset Serial Numbers

Is it possible to add a mechanism to reset the serial numbers?

Mallicious Serial Numbers

Is it possible for byzantine processes to cause the serial numbers (within correct processes) to explode?

If so, is it possible to prevent this?

Thanks for listening!